


Generative Adversarial Networks

Ian Goodfellow, Jean Pouget-Abadie, Mehdi Mirza, Bing Xu, David Warde-Farley, Sherjil Ozair, Aaron Courville, and Yoshua Bengio. Generative adversarial nets. In Advances in Neural Information Processing Systems, pages 2672–2680, 2014.

Comments

Yann LeCun (Facebook AI research director) described it as “the most interesting idea in the last 10 years in Machine Learning.”



What are GANs used for?

The previous deep learning techniques we learned are used for recognition. GAN technique is going to let the machine generate stuff based on what it learned. To note, GAN does not simply memorize the given dataset.

For example, we first learned how to recognize digits like 0-9, then we tried to mimic the shape of digits and created digits in our styles. It is called generation.

Image generation.

Image-to-Image translation

Supervised / unsupervised

Text-to-Image

Super resolution

Photo inpainting

What is a GAN?

Example

How to be a master of producing fake dollars?

Criminor: Produced 1st version fake dollars and the FBI cannot figure them out.

FBI: Found the new fake dollars and successfully figured all 1st version fake dollars out.

Criminor: Produced 2nd version fake dollars and the FBI cannot figure them out.

FBI: Found the new fake dollars and successfully figured all 2nd version fake dollars out.

...

n-th version fake dollars

...

=> Criminor: An expert on making fake dollars

=> FBI: An expert on figuring fake dollars.

In the image generation field

The underlying essence of GAN is simulating the distribution of real data. Then we can sample any random points from this distribution.

If a model only learns the data points rather than the distribution, this model can only memorize the dataset.

A model can generate new data samples when it learns the whole distribution.

Objective for each model

- Discriminator is going to distinguish all fake images.

Intuitively, the objective function is maximizing:

$$J^{(D)} = \text{high_score}(\text{real images}) + \text{low_score}(\text{fake images})$$

- We are familiar with such models, which are going to classify two types of images.

Discriminator is like a binary classifier to distinguish real or fake:

Input a given image, Discriminator function needs to output its probability of reality.

$D(x) = p$. (by softmax or sigmoid)

- Loss function: Cross-entropy (preferred), MSE

[Cross-entropy loss]

Any loss consisting of a negative log-likelihood is a cross-entropy between the empirical distribution defined by the training set and the probability distribution defined by model. For example, mean squared error is the cross-entropy between the empirical distribution and a Gaussian model.

- Cross-entropy loss:

Minimize the divergence between real distribution X_r in the training set and the probability distribution defined by D model.

Maximize the divergence between generated distribution X_g and the probability distribution defined by D model.

Then the loss function for D: θ_D

$$\min_{\theta_D} L^{(D)}(X_r, X_g) = -\mathbb{E}_{x_r \sim X_r} [\log(D(x_r))] - \mathbb{E}_{x_g \sim X_g} [\log(1 - D(x_g))]$$

```
\begin{align*}
\underset{\{\theta_D\}}{\min} L^{\left\{ \left( D \right) \right\}} \left( X_r, X_g \right) = & -\mathbb{E}_{\{x_r \sim X_r\}} \left[ \log \left( D \left( x_r \right) \right) \right] \\
& -\mathbb{E}_{\{x_g \sim X_g\}} \left[ \log \left( 1 - D \left( x_g \right) \right) \right]
\end{align*}
```

- Generator is going to generate real-like images to fool the discriminator.

Intuitively, the objective function is maximizing:

$$J^{(G)} = \text{high_score}(\text{fake images})$$

where high_score is criticized by D , while G can only control the parameters of generating fake images $x_g = G(z)$.

Input a noise vector z to control the generated image $x_g = G(z)$.

Generator function G minimize the divergence between generated distribution X_g and the probability distribution defined by D model.

Then the loss function for G: θ_G

Non-saturating: Min(divergence b/w generated distribution and distribution defined by D model)

$$\min_{\theta_G} L^{(G)}(X_g) = -\mathbb{E}_{x_g \sim X_g} [\log(D(x_g))]$$

```
\underset{\{\theta_G\}}{\min} L^{\left\{ \left( G \right) \right\}} \left( X_g \right) = -\mathbb{E}_{\{x_g \sim X_g\}} \left[ \log \left( D \left( x_g \right) \right) \right]
```

where \textit{D} denotes the discriminator function, \textit{G} denotes the generator function, θ_G is the parameters of the generator, θ_D is the parameters of the discriminator; x_r is sampled from the real distribution X_r , x_g is sampled from the generated distribution X_g , where $x_g = G(z)$ and z is a random noise vector sample from normal distribution $z \sim N(0, I_{\dim(z)})$.

...Beyond the introduction level. Saturating: Max(divergence b/w generated distribution and distribution defined by NOT D model)

$$\mathbb{E}_{x_g \sim X_g} [\log(1 - D(x_g))]$$

Combine these two loss functions together. Minimax problem in Game Theory. (Understand L_D and L_G is more important than this formula)

$$\min_{\theta_G} \max_{\theta_D} L(X_r, X_g) = \mathbb{E}_{x_r \sim X_r} [\log(D(x_r))] + \mathbb{E}_{x_g \sim X_g} [\log(1 - D(x_g))]$$

```
\begin{align*}
&\underset{\theta_G}{\min} \underset{\theta_D}{\max} L(X_r, X_g) = \mathbb{E}_{x_r \sim X_r} [\log(D(x_r))] \\
&+ \mathbb{E}_{x_g \sim X_g} [\log(1 - D(x_g))]
\end{align*}
```

We cannot optimize this combined loss function by changing θ_G and θ_D simultaneously. In practice, we can only train these two models alternatively.

How to train a GAN?

Algorithm

Initialization: random θ_G^0 and θ_D^0

1st iteration

Train D (freeze G)

Input: real images x_r , noise vector $z \sim N(0, 1)$

Calculate: $x_g^0 = G^0(z)$ generated by 0-th version G^0

Output: $D^0(x)$

Optimization: minimize $-\log D^0(x_r) - \log(1 - D^0(x_g^0))$

Update: θ_D^1

Train G (freeze D)

Input: noise vector $z \sim N(0, 1)$

Calculate: $x_g^0 = G^0(z)$ generated by 0-th version G^0


Output: $D^1(G^0(z))$

Optimization: minimize $-\log D^1(G^0(z))$

Update: θ_G^1

... n-th iteration


ReLU and LeakyReLU





ReLU vs LeakyReLU

MLPGAN:

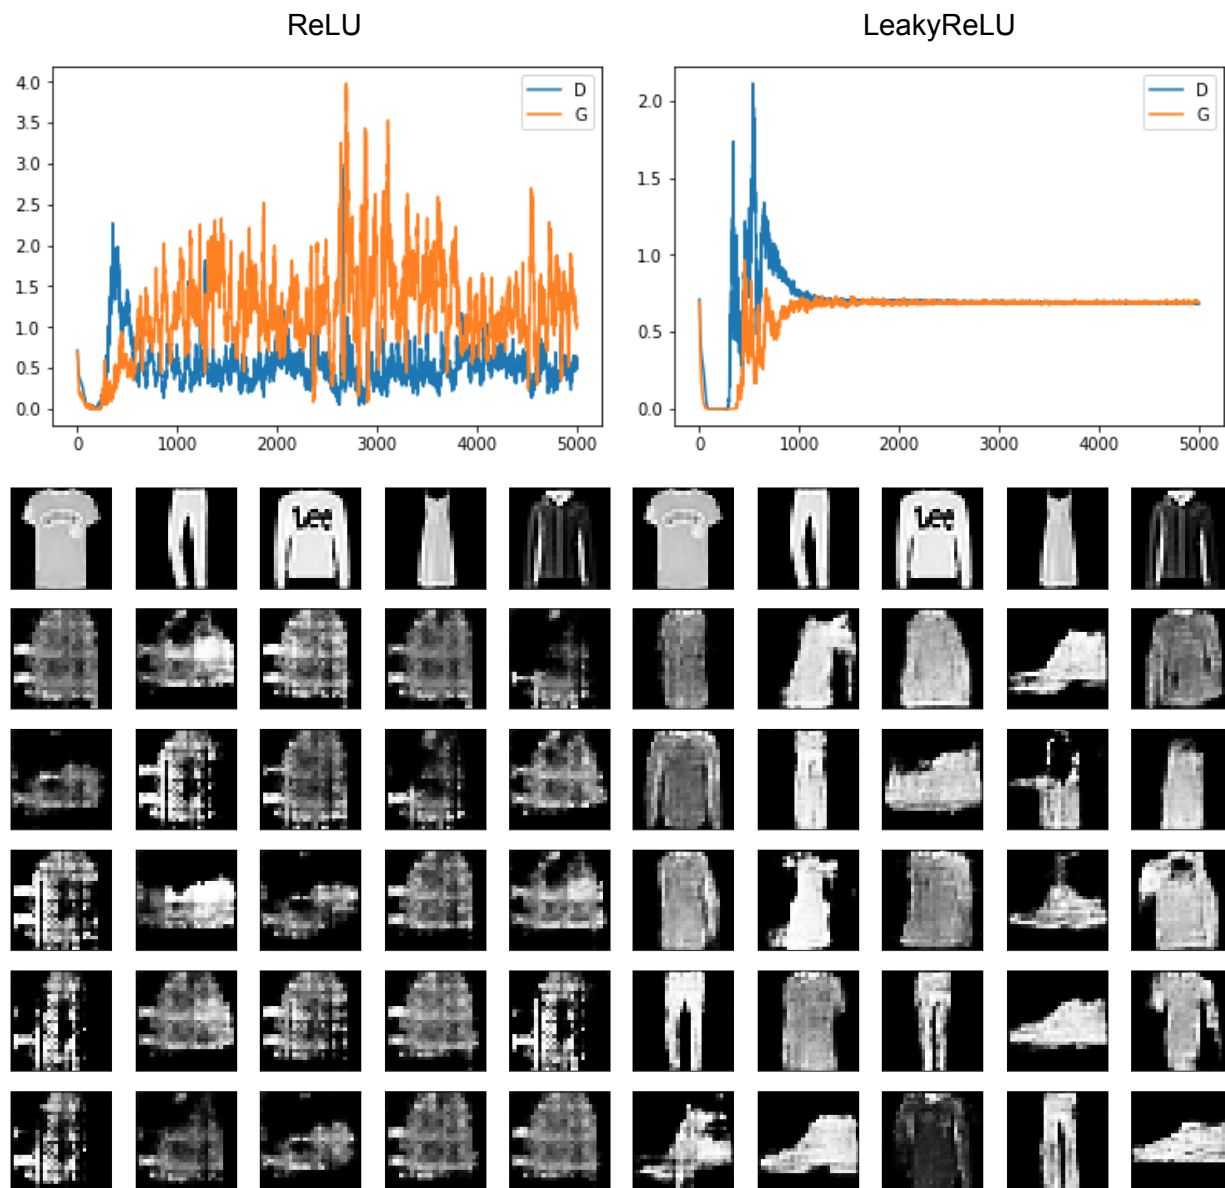
ReLU



LeakyReLU



DCGAN:



The architecture of Generator and Discriminator.

continuous and differentiable functions of D and G → more stable

Discriminator - Downsampling networks

Dense (not suggested)

Fully connected layer is not good at extracting features.

Maxpooling2D (not suggested)

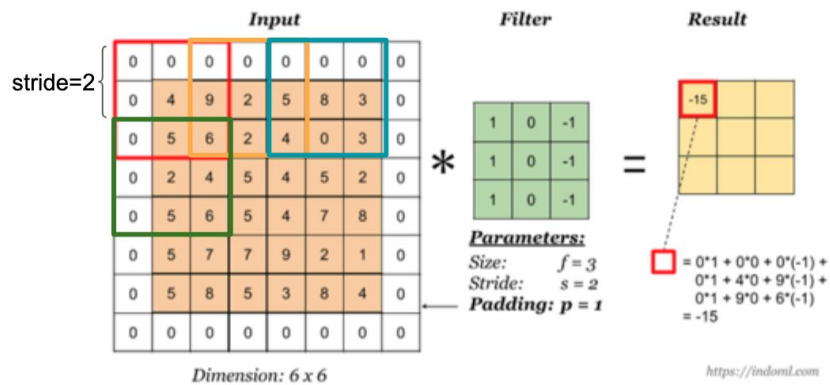
The output is just selecting the maximum input value within the [height, width] window of input values.

There is no weight and no trainable parameters introduced by this operation.

Conv2D(stride=2)

The output is a linear combination of the input values times a weight for each cell in the [height, width] kernel.

These weights become trainable parameters in your model.



Stride is the moving step size.

Generator - Upsampling networks

Dense


Upsampling2D

```
1 [1, 2]
2 Input = [3, 4]
3
4 [1, 1, 2, 2]
5 Output = [1, 1, 2, 2]
6 [3, 3, 4, 4]
7 [3, 3, 4, 4]
```

ConvTranspose / Deconvolution

Stride(2,2)

Stride(1,1), Padding=0



Stride here is the reciprocal of the moving step. E.g. $\text{stride}=2 \Rightarrow \text{moving step}=\frac{1}{2}$. How to move $\frac{1}{2}$ step? Adding zero columns and rows to the original input.

Try our first GAN model with simple MLP-GAN.

Model: 'Discriminator'

Layer (type)	Output Shape	Param #
input_26 (InputLayer)	[(None, 28, 28, 1)]	0
flatten_11 (Flatten)	(None, 784)	0
dense_73 (Dense)	(None, 512)	401920
leaky_re_lu_73 (LeakyReLU)	(None, 512)	0
dense_74 (Dense)	(None, 256)	131328
leaky_re_lu_74 (LeakyReLU)	(None, 256)	0
dense_75 (Dense)	(None, 128)	32896
leaky_re_lu_75 (LeakyReLU)	(None, 128)	0
dense_76 (Dense)	(None, 1)	129
Total params: 566,273		
Trainable params: 566,273		
Non-trainable params: 0		

Model: 'Generator'

Layer (type)	Output Shape	Param #
input_27 (InputLayer)	[(None, 32)]	0
dense_77 (Dense)	(None, 128)	4224
batch_normalization_34 (Batch Normalization)	(None, 128)	512
leaky_re_lu_76 (LeakyReLU)	(None, 128)	0
dense_78 (Dense)	(None, 256)	33024
batch_normalization_35 (Batch Normalization)	(None, 256)	1024
leaky_re_lu_77 (LeakyReLU)	(None, 256)	0
dense_79 (Dense)	(None, 512)	131584
batch_normalization_36 (Batch Normalization)	(None, 512)	2048
leaky_re_lu_78 (LeakyReLU)	(None, 512)	0
dense_80 (Dense)	(None, 784)	402192
reshape_14 (Reshape)	(None, 28, 28, 1)	0
Total params: 574,608		
Trainable params: 572,816		
Non-trainable params: 1,792		

Discriminator with convolution layer

Layer (type)	Output Shape	Param #
input_29 (InputLayer)	[(None, 28, 28, 1)]	0
conv2d_9 (Conv2D)	(None, 14, 14, 32)	320
leaky_re_lu_82 (LeakyReLU)	(None, 14, 14, 32)	0
conv2d_10 (Conv2D)	(None, 7, 7, 64)	18496
leaky_re_lu_83 (LeakyReLU)	(None, 7, 7, 64)	0
conv2d_11 (Conv2D)	(None, 3, 3, 128)	73856
leaky_re_lu_84 (LeakyReLU)	(None, 3, 3, 128)	0
flatten_12 (Flatten)	(None, 1152)	0
dropout_3 (Dropout)	(None, 1152)	0
dense_82 (Dense)	(None, 1)	1153
Total params: 93,825		
Trainable params: 93,825		
Non-trainable params: 0		

Generator with convolution layer

Layer (type)	Output Shape	Param #
input_28 (InputLayer)	[(None, 32)]	0
dense_81 (Dense)	(None, 1152)	38016
leaky_re_lu_79 (LeakyReLU)	(None, 1152)	0
reshape_15 (Reshape)	(None, 3, 3, 128)	0
conv2d_transpose_15 (Conv2DT	(None, 7, 7, 128)	147584
batch_normalization_37 (Batc	(None, 7, 7, 128)	512
leaky_re_lu_80 (LeakyReLU)	(None, 7, 7, 128)	0
conv2d_transpose_16 (Conv2DT	(None, 14, 14, 64)	73792
batch_normalization_38 (Batc	(None, 14, 14, 64)	256
leaky_re_lu_81 (LeakyReLU)	(None, 14, 14, 64)	0
conv2d_transpose_17 (Conv2DT	(None, 28, 28, 1)	577
Total params: 260,737		
Trainable params: 260,353		
Non-trainable params: 384		

Problems in tuning

How to measure the quality of generated images?